

## Adam Marwan

AI Engineer — Remote

Email: adam.marwan11@example.com | Phone: +971594994187

### Professional Summary

Experienced AI Engineer with a strong track record of building production-grade machine learning systems, from data engineering and model training to deployment and monitoring. Skilled in both research prototyping and scalable MLOps practices. Proven ability to lead cross-functional teams and deliver business impact.

### Technical Skills

Data Visualization (Matplotlib/Seaborn), AutoML, Time Series, NLP (BERT, GPT), Keras, Computer Vision (YOLO, Detectron2), Kafka, PyTorch Lightning

### Professional Experience

#### AI Engineer — MachineSense (1 yrs)

- Led an end-to-end image segmentation project for medical imagery; built data pipeline, trained U-Net variants, and reduced labeling time by 40%.
- Optimized model inference using TensorRT and mixed precision; achieved 2.5x throughput improvement on GPU.
- Deployed scalable inference service on AWS with autoscaling, containerization (Docker), and serverless endpoints; reduced latency to <120ms.
- Implemented knowledge distillation to create lightweight transformer models for edge deployment with 3x speedup.

#### Senior AI Engineer — CloudScale (2 yrs)

- Built recommendation system using collaborative filtering + content-based features; increased CTR in A/B test by 12%.
- Optimized model inference using TensorRT and mixed precision; achieved 2.5x throughput improvement on GPU.

#### MLOps Engineer — AIWorks (1 yrs)

- Fine-tuned multilingual speech recognition model (wav2vec2) for domain-specific calls with 5% error reduction.
- Implemented continuous training pipeline using DVC and GitHub Actions to automate model retraining and versioning.

#### AI Engineer — DataForge (2 yrs)

- Designed and implemented an anomaly detection system for fraud detection using autoencoders and XGBoost for post-filtering.
- Implemented knowledge distillation to create lightweight transformer models for edge deployment with 3x speedup.
- Led an end-to-end image segmentation project for medical imagery; built data pipeline, trained U-Net variants, and reduced labeling time by 40%.

### Selected Projects

- Fine-tuned multilingual speech recognition model (wav2vec2) for domain-specific calls with 5% error reduction.
- Optimized model inference using TensorRT and mixed precision; achieved 2.5x throughput improvement on GPU.

### Education

MSc in Data Science — Institute of AI, 2021

### Certifications

- Google Cloud Professional Data Engineer

**Languages**

French (Conversational), English (Fluent)

*Generated sample resume — 2025-09-30*